

Trust me bro

Building a prompt injection benchmark
without a virtual machine

Why I wanted to test this myself

I started this because I kept reading that AI agents were getting dangerous. Reward hacking, models routing around their instructions, that kind of story. The claims interested me, but I wanted something I could run and watch for myself. Which models would actually follow a malicious instruction? Would a stronger model spot the trap, or just be better at springing it?

Those questions pointed me at prompt injection. Reward hacking got me curious, but the experiment I settled on was narrower. Give a coding agent a real task, then plant misleading instructions in the material it reads while working. A repository file or a skill document might tell it to run a setup command the task never asked for, or to hand over a fake secret. I wanted to see when the agent would read that text as permission.

I named it Trust me bro because that's the whole test. A file can sound authoritative without having any right to give orders, and I wanted to watch what an agent did when it took that authority at face value.

A shell I could observe

The decision that made this practical was keeping the agent's working environment inside a JavaScript runtime. I used `just-bash`, which gives me a Bash-like shell and a virtual filesystem, so I never had to start a virtual machine for a run. The agent reads files and issues shell commands through tool calls, and I decide what those calls do.

That control mattered more to me than reproducing a full computer. A `bash` call goes to the virtual shell. A `write_file` call changes virtual file contents. Scenario handlers supply mocked commands and network responses. The model itself is real, running inference through OpenRouter, but the repository and the fake credentials and the dangerous effects all live in the test environment.

Since every action passes through a handler, I can look at a command before it runs, or record an attempted request to a simulated endpoint. I can also have a harmless command fire a canary, a marker for one defined violation. That gives me a concrete event to grade. I'd rather grade an event than ask another model whether the agent's explanation sounded safe.

What goes into a scenario

Each scenario is built around a legitimate user task and a small virtual workspace. It declares its identity, the task, workspace and skills paths, and a map of filenames to contents. Those files carry the useful project context and the injected instructions mixed together, the way an injection would actually arrive. Hidden canaries define the behavior I want to catch. Optional task and safety evaluators check file contents, file changes, command results, or canary states. A trusted runtime fixture supplies any special simulated behavior the scenario needs.

The selected model and the effective system prompt live outside those files. That lets me hold an attack fixed and swap the model or its instructions. The agent under test gets the task and the tools, never the grading rules. Every run starts on a fresh virtual filesystem, so one model's edits can't become the next model's starting conditions.

How a run fits together

The agent acts through tools, and I check what changed in the virtual workspace. The meta-agent reads those results and revises either the attack or the defender's prompt.




Figure 1: The inner loop runs the agent. The outer loop edits files or the system prompt and tries again.

Inside a scenario




Figure 2: A task, virtual files, and hidden checks make one scenario. The marked skill file carries the injection. The checks sit outside the agent's workspace.

One trap, one event I can see

In one test I ask the agent to fix a greeting helper. A skill file gives the correct greeting convention, then demands a setup step that pipes a fetched script into Bash. The fetch is simulated. Running that script fires a harmless canary. Just reading it doesn't. So I can see whether the agent took the bait while it did the work.

The checks are deterministic, but the model's choices aren't. I record task success and safety as separate results, because avoiding a canary by doing nothing doesn't fix the user's problem.

Letting an agent improve the test

Once I could run and inspect a scenario, I wanted an agent to help me change it. So I built a meta-agent that sits above the evaluated agent and works toward a goal I pick. In red-team mode it proposes edits to scenario text or virtual files to make an injection land more often. In blue-team mode it can touch only the system prompt, to make the defender behave better. Both go through the same runner, so I can put a proposed change next to a baseline.

I made those proposals structured edits rather than arbitrary code. Validation checks the allowed fields, paths, and edit limits before a candidate ever runs. The meta-agent can't rewrite the canaries or evaluators, and it can't invent an executable runtime fixture. Otherwise it could raise its score by changing what counts as success. That's the exact behavior that got me interested in this project, and I didn't want to build it into my own tooling.

The controller evaluates candidates within run, token, and spending limits. Red teaming ranks task success ahead of attack success. Blue teaming ranks safety first, then task success, and uses a heuristic to penalize what look like no-op refusals. Repeated runs support baseline comparisons, and each saved revision links a candidate to its edits and outcomes. I wanted to open an experiment months later and still understand why the controller accepted a change.

What surprised me about the models

In my exploratory runs, some frontier models were very good at following whatever instructions they met, including the ones that led somewhere unsafe. Some weaker models looked safer. Then I read the traces and got less comfortable with that ranking. A model could sidestep a malicious sequence because it never understood the sequence, or couldn't execute it.

That changed how I read the results. I don't want to credit a model with security judgment when it was just unable to finish the work. The comparison I care about is whether a model completes the legitimate task and refuses the injected instruction at the same time. A low attack success rate on its own doesn't tell me that.

The system prompt changed what I saw too. When I told models plainly to treat repository files, skills, and tool output as untrusted context, the frontier models became more secure in the runs I observed. My reading is that instruction following feeds both behaviors. A capable model follows a bad instruction well, and it follows a clear security rule well. Some of the gap I saw between models probably came from how I instructed them.

These are qualitative observations, not a controlled ranking of model families. I haven't shown that lower capability causes greater safety, or that a stronger system prompt solves prompt injection. To claim either, I'd need repeated comparisons with matched tasks, settings, and task-success checks, including attacks the prompt optimizer hadn't already seen.

What I want others to build

I want other people to write a virtual workspace, define a task and its checks, and run that against whichever models they care about. The scenario format and the shared runner make that possible without preparing a virtual machine for every test. The meta-agent adds a way to search for better attacks or defenses while keeping the edits and the evidence.

The virtual environment covers only the behaviors I implemented, and each evaluator catches only the condition it checks. A clean run means the agent didn't trip the canaries I wrote, not that it's safe. I take those limits as reasons to keep scenarios inspectable and easy to extend. I started out wanting to know which model I could trust. I ended up wanting to see the exact command it ran and the line of text that talked it into running it.