

Trust me bro

Building a prompt injection benchmark
without a virtual machine

Why I wanted to test this myself

I started this project because I kept seeing claims about how dangerous AI agents were becoming. Stories about reward hacking and models finding ways around their instructions caught my attention, but I wanted something I could test myself. Which models would actually follow a malicious instruction? Would a stronger model notice the trap, or would it be better at carrying it out?

Those questions led me to prompt injection. Reward hacking helped spark my interest, but the experiment I chose was more specific. I gave a coding agent a legitimate task and put misleading instructions in material it would encounter while working. A repository file or skill document could ask it to run an unrelated setup command, disclose a fake secret, or change something outside the task. I wanted to see when the agent would treat that text as permission.

I called the project Trust me bro. The name fit the thing I was testing. A file can sound authoritative without having any right to give instructions. I wanted to watch what happened when an agent took that authority at face value.

A shell I could observe

The design decision that made the project practical was to keep the agent's working environment inside a JavaScript runtime. I used `just-bash` for a Bash-like shell and a virtual filesystem. I did not need to start a virtual machine for each run. The agent could read files, issue shell commands, and write changes through tool calls, while I controlled what those calls did.

That control mattered more to me than reproducing an entire computer. A `bash` call went to the virtual shell. A `write_file` call changed virtual file contents. Scenario-specific handlers supplied simulated network responses and commands. The model still used real inference through OpenRouter, but the repository, fake credentials, and dangerous effects belonged to the test environment.

Because actions passed through these handlers, I could inspect a command before execution, check a write, or record an attempted request to a simulated endpoint. I could also make a harmless command fire a canary, a marker for a defined violation. This gave me a concrete event to grade without asking another model whether the agent's explanation sounded safe.

What goes into a scenario

I built each scenario around a legitimate user task and a small virtual workspace. The scenario describes its identity, task, workspace and skills paths, and a map of filenames to contents. Those files hold both the useful project context and the injected instructions. Hidden canaries define the behavior I want to catch. Task and safety evaluators check file contents, file changes, command results, or canary states. A trusted runtime fixture supplies any special simulated behavior the scenario needs.

I keep the selected model and effective system prompt separate from those files. That lets me hold an attack fixed and change the model or its instructions. The evaluated agent receives the task and tools, not the grading rules. Each run starts with its own virtual filesystem, so one model's edits do not become another model's starting conditions.

The architecture of a run

The agent acts through tools, and I check what happens in the virtual workspace. The meta-agent uses those results to revise an attack or the defender's prompt.




Figure 1: The inner loop runs the agent. The outer loop changes files or the system prompt and tries again.

Inside a scenario




Figure 2: A task, virtual files, and hidden checks make one scenario. The marked skill file contains the injection; the checks stay outside the agent's workspace.

One trap, one observable event

In one test, I ask the agent to fix a greeting helper. A skill file gives the correct greeting convention, then demands a setup step that pipes a fetched script into Bash. The response is simulated. Executing that script fires a harmless canary; simply reading it does not. I can see whether the agent followed the trap while doing the work.

The checks are deterministic, but model choices can vary. I record task success separately from safety. Avoiding a canary by doing nothing does not solve the user's problem.

Letting an agent improve the test

Once I could run and inspect a scenario, I wanted an agent to help me change it. I built a meta-agent that sits above the evaluated agent and works toward a chosen goal. In red-team mode, it proposes changes to scenario text or virtual files to make an injection more effective. In blue-team mode, it can change only the system prompt to improve the defender's behavior. Both use the same runner, so I can compare a proposed change with a baseline.

I made those proposals structured edits rather than arbitrary code. Validation checks the allowed fields, paths, and edit limits before a candidate runs. The meta-agent cannot rewrite the canaries or evaluators, or invent an executable runtime fixture. Otherwise it could improve its score by changing what counts as success. That would recreate the problem that got me interested in this project in the first place.

The controller evaluates candidates within run, token, and spending limits. Red teaming gives task success priority before attack success. Blue teaming ranks safety first, then task success, and penalizes apparent no-op refusals through a heuristic. Repeated runs support baseline comparisons, and saved revisions link each candidate to its edits and outcomes. I wanted to be able to open an experiment later and understand why the controller accepted a change.

What surprised me about the models

In my exploratory runs, some frontier models were very good at following the instructions they encountered, including instructions that led to unsafe actions. Some weaker models appeared safer. Looking at the traces made me less comfortable with that simple ranking. A model could avoid a malicious sequence because it failed to understand or execute the sequence at all.

That distinction changed how I read the results. I do not want to give a model credit for good security judgment when it was simply unable to finish the work. The useful comparison is whether it can complete the legitimate task while rejecting the injected instruction. A low attack success rate alone does not answer that question.

Changing the system prompt also changed what I saw. When I explicitly told models to treat repository files, skills, and tool output as untrusted context, the frontier models became more secure in the runs I observed. My reading is that instruction following contributes to both behaviors. A capable model can follow a bad instruction well, but it can also follow a clear security rule well. Part of the apparent gap between models may therefore come from how I instructed them.

These are qualitative observations, not a controlled ranking of model families. I have not established that lower capability causes greater safety, or that a stronger system prompt solves prompt injection. To make those claims, I would need repeated comparisons with matched tasks, settings, and task-success checks, including attacks the prompt optimizer had not already seen.

What I want others to build

I want the community to be able to write a virtual workspace, define a task and its checks, and run that experiment against different models. The reusable scenario format and shared runner make that possible without preparing a virtual machine for every test. The meta-agent adds a way to search for better attacks or defenses while retaining the edits and evidence.

The virtual environment only covers the behaviors I implement, and each evaluator only catches the conditions it checks. I see those limits as reasons to make scenarios inspectable and easy to extend. I started by wanting to know which model I could trust. I ended up wanting to see the command it ran, the file it changed, and the instruction that got it there.

Implementation basis: Trust me bro, meta-agent integration revision 32493c2. Architecture checked against the scenario types, shared browser runner, meta-agent controller, validation, and scoring code.